

BugOps

Web & AI Engineering

Full-service web development and AI engineering for companies that need both — scoped properly, delivered on timelines we stand behind.

<https://bugops.dev> · hello@bugops.dev

Generated June 21, 2026

HOMEPAGE

Web and AI built to ship — scoped clearly, delivered on time.

BugOps designs and builds production-grade web platforms and AI systems for teams that need both — discovery first, fixed scope before build.

Focus areas: Scoped MVPs · Production AI · Integrations

Primary CTA: Book discovery · Secondary: Browse services

Trusted by companies across 12+ industries worldwide

COMPANY STATS

25+ Skilled professionals

Engineers, designers, and AI specialists with deep domain expertise across web and machine learning.

6+ Years of experience

Combined industry experience delivering production systems for startups, SMBs, and enterprise teams.

100+ Projects delivered

Web platforms, AI agents, analytics pipelines, and custom applications shipped across 12+ industries.

98% Client satisfaction

Long-term partnerships built on clear scoping, transparent delivery, and systems that stay running.

INDUSTRIES

Industries

Domain-aware web development and AI engineering — built around your compliance, data, and user workflows.

EdTech

Learning platforms, student portals, and AI tutoring systems that personalize education at scale.

HealthTech

Patient portals, clinical dashboards, and HIPAA-aware AI for documentation and care coordination.

FinTech & banking

Secure web platforms, fraud detection, and AI-driven analytics for financial services.

Retail & e-commerce

High-conversion storefronts, inventory systems, and recommendation engines for online retail.

B2B SaaS

Multi-tenant platforms, customer portals, and embedded AI features for software products.

Real estate

Property platforms, CRM integrations, and AI tools for listings, valuations, and lead management.

Legal & compliance

Document intelligence, contract review workflows, and secure client portals for legal teams.

Manufacturing

Supply chain dashboards, quality inspection with computer vision, and operational analytics.

Hospitality & travel

Booking platforms, itinerary tools, and AI assistants for guest experience and operations.

Media & entertainment

Content platforms, recommendation systems, and AI-powered production and distribution tools.

Government & public sector

Accessible citizen portals, secure internal systems, and document processing at scale.

Energy & utilities

Field operations dashboards, predictive maintenance models, and asset monitoring platforms.

SERVICES OVERVIEW

Web platforms and AI systems scoped to ship — not slide decks

Pick what you need — or book discovery and we'll map the right mix of web, AI, and delivery support.

WEB DEVELOPMENT

Web development

Sites and platforms built to convert, scale, and stay maintainable

From marketing sites to full-scale platforms — built for performance, security, and the workflows your team runs every day.

Corporate & business websites

Business websites — Marketing sites that convert — and your team can actually update

Custom, mobile-first sites with CMS, booking, and lead capture — designed to convert and easy for your team to maintain.

- Custom UI/UX design
- Mobile-first development
- CMS & content workflows
- SEO & launch support

Web applications & portals

Web platforms — Portals and dashboards your users rely on every day

Customer portals, admin dashboards, and internal tools with authentication, APIs, and real-time features.

- Full-stack development
- Database architecture
- Third-party integrations
- Admin dashboards

Enterprise web systems

Enterprise systems — High-traffic platforms with security and architecture that scale

High-traffic platforms with strict security, compliance requirements, and architecture built to scale.

- System architecture
- Security & compliance

- Performance optimization
- DevOps & CI/CD

AI SOLUTIONS

AI solutions

Production AI wired into your stack — with guardrails, not gimmicks

Production-grade AI — agents, chatbots, analytics, and custom applications integrated with your existing stack and data.

AI agent development

Agents that run workflows in production — not just in demos

Autonomous workflows with guardrails

Multi-step agents that read context, make decisions within policy, and execute across your CRM, email, and internal tools.

Use cases

- Lead qualification & routing
- Operations automation
- Document processing pipelines

Capabilities

- Agent orchestration
- Tool & API integration
- Human-in-the-loop review

RAG chatbot development

Answers grounded in your data — with escalation when it matters

Answers grounded in your data

Retrieval-augmented chatbots trained on your docs, tickets, and product content — with escalation when confidence is low.

Use cases

- Customer support
- Sales qualification
- Internal knowledge bases

Capabilities

- RAG architecture
- Vector search
- Source citations & logging

AI analytics & insights

From raw data to decisions — without rebuilding spreadsheets every week

From raw data to decisions

Automated trend detection, anomaly alerts, and executive dashboards — replacing manual spreadsheet cycles.

Use cases

- Operational reporting

- Trend & anomaly detection
- Executive dashboards

Capabilities

- Data pipelines
- Automated analysis
- Real-time visualizations

Custom AI applications

Purpose-built AI products — scoped, shipped, and owned by you

Purpose-built AI products

End-to-end AI products combining LLMs, retrieval, and your business logic for industry-specific problems.

Use cases

- Vertical SaaS features
- Internal AI platforms
- Customer-facing AI tools

Capabilities

- Product design & UX
- Custom model orchestration
- Production deployment

AI integration & automation

AI inside the tools you already use — without replacing your stack

AI inside your existing stack

Embed AI into CRMs, ERPs, and custom software — automating repetitive work without replacing systems you rely on.

Use cases

- CRM enrichment
- Workflow automation
- API-driven AI features

Capabilities

- System integration
- Event-driven automation
- Monitoring & fallbacks

NLP & document intelligence

Structure from unstructured documents — at the volume you actually receive

Extract structure from unstructured data

Classify, summarize, and extract fields from contracts, invoices, forms, and correspondence at scale.

Use cases

- Contract review
- Invoice processing
- Compliance document intake

Capabilities

- Document parsing
- Entity extraction

- Classification pipelines

Computer vision solutions

Visual data processed automatically — with models tuned for your environment

Visual data, automated

Image and video analysis for quality control, asset tagging, identity verification, and visual search.

Use cases

- Quality inspection
- Visual search
- Identity verification

Capabilities

- Model training & tuning
- Edge & cloud deployment
- Real-time inference

Predictive AI & ML models

Forecasts and scores you can act on — not models stuck in a notebook

Forecast what happens next

Custom models for demand forecasting, churn prediction, risk scoring, and recommendation engines.

Use cases

- Demand forecasting
- Churn & risk scoring
- Recommendation systems

Capabilities

- Feature engineering
- Model training & validation
- MLOps & retraining

HOW WE DELIVER

How We Deliver

The work that happens around every build — before, during, and after the code.

Product Discovery

Scope Before Quotes — Before any code gets written, we map your goals, users, and constraints in a structured discovery phase — so the build that follows is scoped against a real plan, not a guess.

- Stakeholder/goal mapping
- User & competitor research
- Technical feasibility assessment
- Scoped deliverable & timeline document

MVP Development

From Idea to Working Product, Fast — We build only what's essential to test your idea in the real world — functional, scalable MVPs for founders who need to validate before they build everything.

- Feature prioritization (build only what's essential)
- Rapid build sprints

- Built-in scalability path for post-validation growth
- Founder-friendly fixed-scope pricing

UI/UX Design

Interfaces People Actually Want to Use — Clean, intuitive design across web and AI products — built around real user flows, not just visual polish, so what we design is what actually converts.

- User flow mapping
- Wireframes & high-fidelity design
- Design systems for consistency at scale
- Usability-focused iteration, not just visual polish

Vibe Coding Cleanup

Fix What AI Tools Started — Built something fast with Lovable, Replit, or Cursor and now it's fragile under real use? We audit, refactor, and rebuild AI-generated codebases into something production-stable — without starting over.

- Full codebase audit
- Security & scalability risk assessment
- Refactor of AI-generated code into maintainable structure
- Documentation handoff so the codebase isn't a black box anymore

Digital Transformation

Modernize Without the Disruption — For teams running on legacy systems or manual processes — we map a practical path to modern infrastructure and AI-assisted workflows, without breaking what already works.

- Legacy system/process audit
- Modernization roadmap (no full rip-and-replace required)
- AI-assisted workflow integration
- Phased rollout to avoid operational disruption

PROCESS

Discovery before dates. Scope before quotes.

Fixed-scope MVPs for validation, or ongoing builds for teams that need a long-term engineering partner — scoped in discovery, not guessed from a form.

Step 01

Discovery

Understand the problem. We map your goals, users, data, and constraints before proposing any solution or timeline.

Step 02

Scoping

Document everything. Requirements, architecture, deliverables, and effort — written and agreed before development starts.

Step 03

Build

Ship in stages. Iterative delivery with staging reviews, QA, and regular check-ins so nothing surprises you at launch.

Step 04

Launch & support

Stay running. Deployment, monitoring, handoff documentation, and post-launch support through go-live and beyond.

ENGAGEMENT MODELS

Fixed-scope MVP

Defined deliverables, timeline, and budget — built to validate before you invest in the full product.

Ongoing build

Iterative delivery with staging reviews and a team that stays on through launch and beyond.

TECHNOLOGY STACK

Modern frameworks and platforms we use to ship web platforms, AI systems, and production infrastructure.

Web

Next.js — Marketing sites, customer portals & full-stack web applications

React — Interactive dashboards, admin tools & component-driven UIs

TypeScript — Type-safe frontends, APIs & shared code across the stack

Node.js — REST & real-time APIs, integrations & backend services

Tailwind CSS — Design systems, responsive UI & rapid product iteration

Vercel — Edge deployment, previews & production web hosting

Docker — Containerized deployment, CI/CD & scalable production ops

Prisma — Type-safe database access & schema-driven backends

GraphQL — Flexible APIs for portals, mobile clients & integrations

Supabase — Auth, storage & realtime features for modern web apps

Kubernetes — Enterprise orchestration, scaling & resilient deployments

Google Cloud — Cloud infrastructure, managed AI services & enterprise hosting

AI

Python — ML pipelines, data processing & predictive models

OpenAI — LLM chatbots, agents & document intelligence

LangChain — AI agents, RAG orchestration & tool integrations

Hugging Face — Model hosting, fine-tuning & NLP experimentation

PyTorch — Custom model training, computer vision & deep learning

TensorFlow — Production ML models & large-scale inference pipelines

Anthropic — Enterprise-grade LLM assistants & reasoning workflows

Pinecone — Vector search for grounded RAG chatbots & knowledge bases

FastAPI — High-performance AI APIs & model serving endpoints

NVIDIA — GPU-accelerated training & computer vision workloads

Data

PostgreSQL — Transactional data, analytics stores & RAG metadata

Redis — Caching, sessions & real-time job queues

MongoDB — Flexible document storage for product & workflow data

DEMO AI VS PRODUCTION AI

Demo AI vs production AI

Most AI projects stall between a working prototype and something your team can run every day. We build for the second.

Demo agent

Clean input, happy path

- Works on well-structured test data — impressive in a meeting, fragile in the wild
- Single linear path with no recovery when something unexpected happens
- Prompts tuned for ideal inputs, not ambiguous or incomplete messages
- Token usage unchecked — fine for a demo, unsustainable at scale
- Failures go unnoticed until a user reports them

Production agent

How we build at BugOps

- Input validation and guardrails before anything reaches the model
- Branching logic, retries, and graceful escalation — not silent failures
- Human-in-the-loop when confidence is low, with full context preserved
- Token budgets, rate limits, and cost monitoring from day one
- Logging, evals, and monitoring so issues surface before users do

PORTFOLIO / CASE STUDIES

SmartSupport

B2B SaaS · Customer support · Portfolio project

Tech: Next.js, LangChain, Pinecone, Node.js, PostgreSQL

Key metric: 68% faster ticket resolution

Support team was drowning in triage. We built a dashboard where AI drafts replies from past tickets and docs — agents review before anything goes out.

Challenge

Every ticket was read, tagged, and routed by hand. Volume kept climbing and senior people spent most of their day sorting, not fixing.

Web layer

Internal dashboard for the support team: queue view, suggested reply, one-click edit and send. Human approval on every outbound message.

AI layer

RAG over resolved tickets and product docs. Classifies severity and drafts a grounded reply. Falls back to a person when retrieval is weak.

Data layer

Ticket history feeds back into the vector store as cases close, so retrieval improves without a separate retraining cycle.

Impact

- 68% reduction in average resolution time
- Severity mis-tags under 5% of tickets
- Senior staff moved off triage to escalations only

SmartTrip

Travel · Consumer product · Portfolio project

Tech: Next.js, OpenAI, PostgreSQL, Redis

Key metric: Itineraries in under 8 seconds

Client wanted personalized trip plans, not another listicle. We shipped a web app that turns dates, budget, and interests into a day-by-day itinerary.

Challenge

Static content was not converting. Users wanted a plan they could save and adjust — generated in seconds on mobile.

Web layer

Single-page app with optimistic updates so generation feels instant. Mobile-first — most planning happens on a phone.

AI layer

Structured inputs in, structured day-by-day plan out. Prompting tuned for geographic coherence and realistic pacing.

Data layer

User profiles, saved itineraries, and a curated POI dataset so suggestions stay tied to real locations.

Impact

- Average generation under 8 seconds
- Higher save rate vs. the old static content format

CollaboraX

B2B · Team collaboration · Portfolio project

Tech: React, Node.js, PostgreSQL, WebSockets

Key metric: Zero isolation incidents post-launch

Multi-tenant workspace for distributed teams — real-time editing without the sync bugs that usually show up in collaborative tools.

Challenge

Simultaneous edits kept causing desync and data leaks between client workspaces. They needed strict isolation and predictable state.

Web layer

Centralized state management for concurrent edits — no race conditions, no UI drifting from the server.

AI layer

Not part of this engagement — web and data only.

Data layer

Multi-tenant schema with workspace-level isolation. WebSocket layer tuned for low latency under concurrent sessions.

Impact

- Stable real-time performance under heavy concurrent use
- No reported cross-workspace data leaks after launch

Trendforce

Operations · Leadership reporting · Portfolio project

Tech: Python, PostgreSQL, LangChain, Custom dashboards

Key metric: Reporting cut from days to hours

Leadership was waiting on manual spreadsheet pulls. We wired ingestion pipelines and dashboards that flag shifts without someone rebuilding the sheet each week.

Challenge

Operational data lived in disconnected spreadsheets. Spotting a trend meant a multi-day manual compile — and the numbers were often stale by the time they landed.

Web layer

Lightweight dashboards for leadership — not a full product, but something they could open without asking ops for a export.

AI layer

Automated trend and anomaly detection on ingested metrics — replaces hand-built month-over-month comparisons.

Data layer

Pipelines consolidate sources into one structured store with the formulas leadership already trusted.

Impact

- Multi-day reporting cycles replaced with always-current views
- Trends surfaced without waiting on a manual pull

MISSION & VISION

Your long-term web and AI technology partner

We partner with organizations to design, build, and maintain web platforms and AI systems — from first scope to production and beyond.

Our mission

Deliver production-grade web development and AI engineering that drives measurable business outcomes — with clear scoping, honest timelines, and no vendor handoffs between your site and your AI layer.

Our vision

To be the engineering partner teams trust for custom web platforms and applied AI — known for shipping systems that work in production, not just in demos.

TEAM

Led by co-founders with deep expertise in full-stack web engineering and applied AI — backed by a senior engineering team across design, development, and machine learning.

Irfan Ali — Co-Founder

Mian Asadullah — Co-Founder

TRUST & PARTNERS

Trusted companies referenced: Linear, Notion, Airtable, Intercom, Zendesk, Miro, Zapier, HubSpot

Technology partners: OpenAI, Claude, Anthropic, Google Cloud, LangChain, Hugging Face, NVIDIA, Mistral AI

BOOK A CALL

Tell us what you are building — we will map scope and fit.

We ask about goals, timeline, and whether we are the right team. Fixed-scope MVP or ongoing build — mapped in discovery, not on a sales script.

What to expect: Listen first, scope second, quote third.

We reply within one business day.

HOMEPAGE FAQ

Web development

What types of websites and platforms do you build?

Corporate and marketing sites, customer portals, admin dashboards, and enterprise web systems — from CMS-driven business sites to full-stack applications with auth, APIs, and real-time features.

Do you handle design, development, and launch?

Yes. We deliver custom UI/UX, mobile-first development, CMS workflows, SEO setup, and go-live support — so your team can maintain content after launch without starting over.

Can you scale a site from small business to enterprise?

We scope in tiers: business websites, web applications, and enterprise systems. You can start with a focused build and expand architecture, security, and DevOps as traffic and compliance requirements grow.

AI solutions

Which AI services do you offer?

AI agents, RAG chatbots, analytics and insights, custom AI applications, system integration and automation, NLP and document intelligence, computer vision, and predictive ML — all built for production, not demos.

How do you integrate AI with our existing stack?

We connect to your CRM, ERP, databases, and internal tools via APIs and event-driven workflows — with guardrails, logging, human-in-the-loop review, and fallbacks when models are uncertain.

Do you build RAG chatbots on our own data?

Yes. We design retrieval-augmented chatbots grounded in your docs, tickets, and product content — with vector search, source citations, and escalation paths when confidence is low.

Can you ship AI agents for internal operations?

We build multi-step agents for lead routing, document processing, support triage, and ops automation — with policy boundaries and tool access scoped to what your team actually needs.

Working together

Do you handle both web development and AI under one team?

Yes. BugOps is a full-service web and AI engineering team — one scope, one timeline, and one team responsible for your platform and intelligence layer.

Which industries do you work with?

We serve EdTech, HealthTech, FinTech, retail, B2B SaaS, real estate, legal, manufacturing, hospitality, media, government, energy, and more — with domain-aware compliance and data workflows.

How does pricing and scoping work?

Every project starts with discovery. We document requirements first, then provide a fixed quote — no guessing from a contact form alone.

How do we get started?

Book an intro call or send a message through our contact form. We will discuss your requirements, timeline, and fit. If we are not the right partner, we will tell you directly.

WEB DEVELOPMENT FAQ

What types of websites and platforms do you build?

Corporate and marketing sites, customer portals, admin dashboards, and enterprise web systems — from CMS-driven business sites to full-stack applications with auth, APIs, and real-time features.

How long does a typical engagement take?

Focused MVPs often run 4–12 weeks. Mid-complexity web or AI builds run 3–6 months. We scope timeline in discovery — written and agreed before development starts.

Who owns the IP and source code?

You do. Code is developed in your repositories where possible, with documentation and handoff included. We assign full IP ownership on signing.

Can you integrate with our existing tools and data?

Yes. We connect via APIs to CRMs, ERPs, databases, and internal systems — with scoped tool access, logging, and fallbacks when integrations fail.

Do you work with regulated or sensitive data?

We configure projects for GDPR, HIPAA-aware workflows, and sector-specific requirements as part of discovery. Sensitive workloads can use private cloud endpoints or self-hosted models when needed.

What happens after launch?

Every project ships with monitoring and handoff documentation. Ongoing support covers performance, security patches, and feature work — with the same engineers who built it.

AI SERVICE FAQS

AI agent development

What makes your agents production-ready?

Policy boundaries, tool access scoped to the workflow, human-in-the-loop escalation, cost controls, and monitoring — not just a prompt that works on demo data.

How long does a typical engagement take?

Focused MVPs often run 4–12 weeks. Mid-complexity web or AI builds run 3–6 months. We scope timeline in discovery — written and agreed before development starts.

Who owns the IP and source code?

You do. Code is developed in your repositories where possible, with documentation and handoff included. We assign full IP ownership on signing.

Can you integrate with our existing tools and data?

Yes. We connect via APIs to CRMs, ERPs, databases, and internal systems — with scoped tool access, logging, and fallbacks when integrations fail.

Do you work with regulated or sensitive data?

We configure projects for GDPR, HIPAA-aware workflows, and sector-specific requirements as part of discovery. Sensitive workloads can use private cloud endpoints or self-hosted models when needed.

What happens after launch?

Every project ships with monitoring and handoff documentation. Ongoing support covers performance, security patches, and feature work — with the same engineers who built it.

RAG chatbot development

How do you keep chatbot answers grounded?

Retrieval-augmented generation over your docs and tickets, source citations, confidence thresholds, and escalation to a person when retrieval is weak.

How long does a typical engagement take?

Focused MVPs often run 4–12 weeks. Mid-complexity web or AI builds run 3–6 months. We scope timeline in discovery — written and agreed before development starts.

Who owns the IP and source code?

You do. Code is developed in your repositories where possible, with documentation and handoff included. We assign full IP ownership on signing.

Can you integrate with our existing tools and data?

Yes. We connect via APIs to CRMs, ERPs, databases, and internal systems — with scoped tool access, logging, and fallbacks when integrations fail.

Do you work with regulated or sensitive data?

We configure projects for GDPR, HIPAA-aware workflows, and sector-specific requirements as part of discovery. Sensitive workloads can use private cloud endpoints or self-hosted models when needed.

What happens after launch?

Every project ships with monitoring and handoff documentation. Ongoing support covers performance, security patches, and feature work — with the same engineers who built it.

AI analytics & insights

Do you replace our existing reporting tools?

Usually we consolidate data into pipelines and dashboards your team already trusts — adding automated trend and anomaly detection, not another siloed tool.

How long does a typical engagement take?

Focused MVPs often run 4–12 weeks. Mid-complexity web or AI builds run 3–6 months. We scope timeline in discovery — written and agreed before development starts.

Who owns the IP and source code?

You do. Code is developed in your repositories where possible, with documentation and handoff included. We assign full IP ownership on signing.

Can you integrate with our existing tools and data?

Yes. We connect via APIs to CRMs, ERPs, databases, and internal systems — with scoped tool access, logging, and fallbacks when integrations fail.

Do you work with regulated or sensitive data?

We configure projects for GDPR, HIPAA-aware workflows, and sector-specific requirements as part of discovery. Sensitive workloads can use private cloud endpoints or self-hosted models when needed.

What happens after launch?

Every project ships with monitoring and handoff documentation. Ongoing support covers performance, security patches, and feature work — with the same engineers who built it.

Custom AI applications

Can you build a full AI product, not just a feature?

Yes — end-to-end products combining UX, model orchestration, data layers, and production deployment for customer-facing or internal platforms.

How long does a typical engagement take?

Focused MVPs often run 4–12 weeks. Mid-complexity web or AI builds run 3–6 months. We scope timeline in discovery — written and agreed before development starts.

Who owns the IP and source code?

You do. Code is developed in your repositories where possible, with documentation and handoff included. We assign full IP ownership on signing.

Can you integrate with our existing tools and data?

Yes. We connect via APIs to CRMs, ERPs, databases, and internal systems — with scoped tool access, logging, and fallbacks when integrations fail.

Do you work with regulated or sensitive data?

We configure projects for GDPR, HIPAA-aware workflows, and sector-specific requirements as part of discovery. Sensitive workloads can use private cloud endpoints or self-hosted models when needed.

What happens after launch?

Every project ships with monitoring and handoff documentation. Ongoing support covers performance, security patches, and feature work — with the same engineers who built it.

AI integration & automation

Will AI integration replace our current systems?

No. We embed AI into CRMs, ERPs, and custom software — automating repetitive work while keeping the systems your team relies on.

How long does a typical engagement take?

Focused MVPs often run 4–12 weeks. Mid-complexity web or AI builds run 3–6 months. We scope timeline in discovery — written and agreed before development starts.

Who owns the IP and source code?

You do. Code is developed in your repositories where possible, with documentation and handoff included. We assign full IP ownership on signing.

Can you integrate with our existing tools and data?

Yes. We connect via APIs to CRMs, ERPs, databases, and internal systems — with scoped tool access, logging, and fallbacks when integrations fail.

Do you work with regulated or sensitive data?

We configure projects for GDPR, HIPAA-aware workflows, and sector-specific requirements as part of discovery. Sensitive workloads can use private cloud endpoints or self-hosted models when needed.

What happens after launch?

Every project ships with monitoring and handoff documentation. Ongoing support covers performance, security patches, and feature work — with the same engineers who built it.

NLP & document intelligence

What document types can you process?

Contracts, invoices, forms, correspondence, and compliance intake — with classification, extraction, and routing into your existing workflows.

How long does a typical engagement take?

Focused MVPs often run 4–12 weeks. Mid-complexity web or AI builds run 3–6 months. We scope timeline in discovery — written and agreed before development starts.

Who owns the IP and source code?

You do. Code is developed in your repositories where possible, with documentation and handoff included. We assign full IP ownership on signing.

Can you integrate with our existing tools and data?

Yes. We connect via APIs to CRMs, ERPs, databases, and internal systems — with scoped tool access, logging, and fallbacks when integrations fail.

Do you work with regulated or sensitive data?

We configure projects for GDPR, HIPAA-aware workflows, and sector-specific requirements as part of discovery. Sensitive workloads can use private cloud endpoints or self-hosted models when needed.

What happens after launch?

Every project ships with monitoring and handoff documentation. Ongoing support covers performance, security patches, and feature work — with the same engineers who built it.

Computer vision solutions

Do you deploy on edge or cloud?

Both — depending on latency, privacy, and hardware constraints. We train, tune, and deploy for real-time inference where your use case requires it.

How long does a typical engagement take?

Focused MVPs often run 4–12 weeks. Mid-complexity web or AI builds run 3–6 months. We scope timeline in discovery — written and agreed before development starts.

Who owns the IP and source code?

You do. Code is developed in your repositories where possible, with documentation and handoff included. We assign full IP ownership on signing.

Can you integrate with our existing tools and data?

Yes. We connect via APIs to CRMs, ERPs, databases, and internal systems — with scoped tool access, logging, and fallbacks when integrations fail.

Do you work with regulated or sensitive data?

We configure projects for GDPR, HIPAA-aware workflows, and sector-specific requirements as part of discovery. Sensitive workloads can use private cloud endpoints or self-hosted models when needed.

What happens after launch?

Every project ships with monitoring and handoff documentation. Ongoing support covers performance, security patches, and feature work — with the same engineers who built it.

Predictive AI & ML models

How do you validate models before production?

Feature engineering, holdout validation, and MLOps for retraining — with monitoring for drift so models stay useful after launch.

How long does a typical engagement take?

Focused MVPs often run 4–12 weeks. Mid-complexity web or AI builds run 3–6 months. We scope timeline in discovery — written and agreed before development starts.

Who owns the IP and source code?

You do. Code is developed in your repositories where possible, with documentation and handoff included. We assign full IP ownership on signing.

Can you integrate with our existing tools and data?

Yes. We connect via APIs to CRMs, ERPs, databases, and internal systems — with scoped tool access, logging, and fallbacks when integrations fail.

Do you work with regulated or sensitive data?

We configure projects for GDPR, HIPAA-aware workflows, and sector-specific requirements as part of discovery. Sensitive workloads can use private cloud endpoints or self-hosted models when needed.

What happens after launch?

Every project ships with monitoring and handoff documentation. Ongoing support covers performance, security patches, and feature work — with the same engineers who built it.

CONTACT

Tell us about your project — web, AI, or both. We will respond with next steps and an honest fit assessment.

Email: hello@bugops.dev

Phone: +1 (555) 000-0000

We reply within one business day.

FOOTER

Ready to start your project? Web platforms, AI systems, or both — scoped properly from day one.

© 2026 BugOps. All rights reserved. Web & AI Engineering